Relational parametricity and "theorems for free"

A tutorial, with example code in Scala

Sergei Winitzki

Academy by the Bay

2021-09-04

Motivation for parametricity. "Theorems for free"

Parametricity: all fully parametric functions satisfy their naturality laws

- Naturality law: code must work in the same way for all types

```
def headOption[A]: List[A] => Option[A] = {
  case Nil              => None
  case head :: tail     => Some(head)
}
```

- **"Fully parametric"** code: use only type parameters, no JVM reflection
  - Naturality laws are "for free" only if all code is fully parametric
- Naturality law for `headOption`: for any `x: List[A]` and `f: A => B`,

```
headOption(x).map(f) == headOption(x.map(f))
```

Parametricity theorems work only if the code is "fully parametric"
Parametricity theorems apply only to a subset of a programming language

- Usually, it is a certain flavor of typed lambda calculus

Examples of code that fails parametricity

Explicit matching on type parameters using JVM reflection:

```
def badHeadOpt[A]: List[A] => Option[A] = {
  case Nil                    => None
  case (head: Int) :: tail => None
  case head :: tail           => Some(head)
}
```

Using typeclasses: define typeclass NotInt[A] returning true unless A = Int

```
def badHeadOpt[A]: List[A] => Option[A] = {
  case h :: tail if implicitly[NotInt[A]]()   => Some(h)
  case _ => None
}
```

Failure of naturality law:

```
scala> badHeadOpt(List(10, 20, 30).map(x => s"x = $x"))
res0: Option[String] = Some(x = 10)

scala> badHeadOpt(List(10, 20, 30)).map(x => s"x = $x")
res1: Option[String] = None
```

Fully parametric programs are written using the 9 code constructions:

```
def fmap[A, B](f: A => B): List[(A, A)] => List[(B, B)] = { // 3
   case Nil             => Nil
//   8   1                  1,7
   case head :: tail    => (f (head._1), f (head._2)) :: fmap(f)(tail)
//   8       6               2 4     6  5 2 4     6    7   9
}           // This code uses each of the nine allowed constructions.
```

1. Use `Unit` value (or equivalent type), e.g. `()`, `Nil`, `None`
2. Use bound variable (a given argument of the function)
3. Create a function: `{ x => expr(x) }`
4. Use a function: `f(x)`
5. Create a product: `(a, b)`
6. Use a product: `p._1` (or via pattern matching)
7. Create a co-product: `Left[A, B](x)`
8. Use a co-product: `{ case ... => ... }` (pattern matching)
9. Use a recursive call: e.g., `fmap(f)(tail)` within the code of `fmap`

Why we need relational parametricity

"Relational parametricity" is a method for proving parametricity theorems

- Main papers: Reynolds (1983) and Wadler "Theorems for free" (1989)
  - Those papers are a bit outdated and also hard to understand
- There are very few pedagogical tutorials on relational parametricity
  - "On a relation of functions" by R. Backhouse (1990)
  - "The algebra of programming" by R. Bird and O. de Moor (1997)

This tutorial does *not* follow any of the above but derives equivalent results

- Alternative approach: prove "dinaturality" (de Lataillade, Voigtländer)
  - Dinaturality is a consequence of relational parametricity
  - In practice, dinaturality laws are sufficient in most cases
  - But some proofs still need full relational parametricity

Motivating relational parametricity. I. Naturality laws

Naturality law: applying `t[A]: F[A] => G[A]` *before* `_.map(f)` equals applying `t[B]: F[B] => G[B]` *after* `_.map(f)` for any function `f: A => B`
Naturality laws need *lifting* `f: A => B` to `F[A] => F[B]` and `G[A] => G[B]`

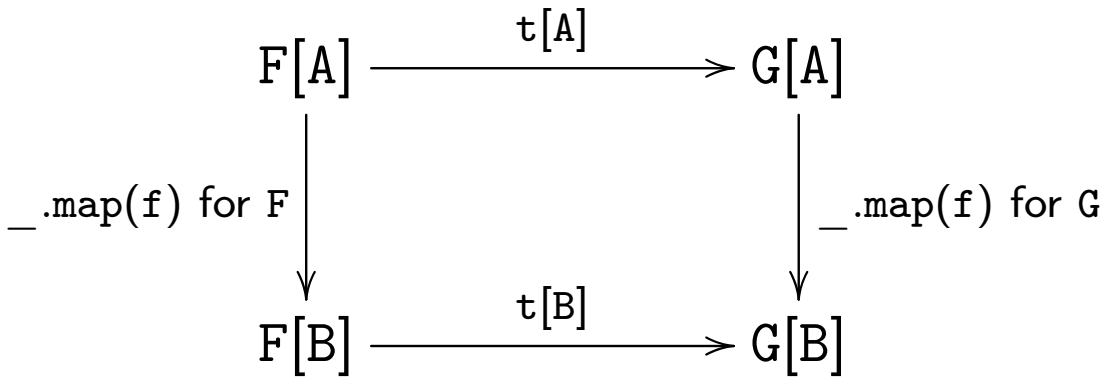

- Proof of the naturality law requires induction on the code of `t[A]`
  - This code is built up by combining the 9 code constructions
  - This code may include sub-expressions of types not covariant in `A`

Motivating relational parametricity. II. The difficulty

Cannot lift `f: A => B` to `F[A] => F[B]` when `F[_]` is not covariant!

- For covariant `F[_]` we lift `f: A => B` to `fmap(f): F[A] => F[B]`
- For contravariant `F[_]` we lift `f: B => A` to `cmap(f): F[A] => F[B]`

In general, `F[_]` will be neither covariant nor contravariant

- Example: `foldLeft` with respect to type parameter `A`

  ```
  def foldLeft[T, A]: List[T] => (T => A => A) => A => A
  ```

- This is *not* of the form `F[A] => G[A]` with covariant `F[_]` and `G[_]`
  - Some occurrences of `A` are in covariant positions but other occurrences are in contravariant positions, all mixed up

Motivating relational parametricity. III. Liftings

The solution involves three nontrivial steps:

1. Replace functions `f: A => B` by relations `r: A <=> B`
   - Instead of `b == f(a)`, we will write: `(a, b) in r`
2. Turns out, we can lift `r: A <=> B` to `rmap(r): F[A] <=> F[B]`
3. Reformulate the naturality law of `t` via relations: for any `r: A <=> B`,

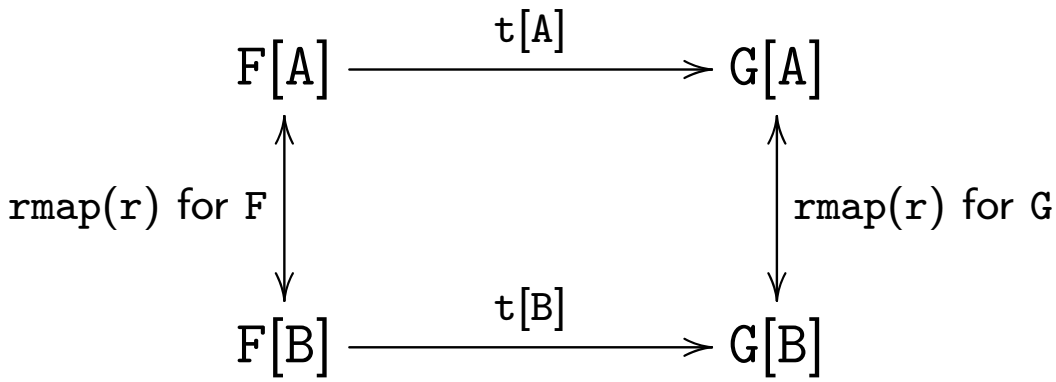

To read the diagram: the starting values are on the left
For any `r: A <=> B`, for any `fa: F[A]` and `fb: F[B]` such that `(fa, fb) in rmap_F(r)`, we require `(t(fa), t(fb)) in rmap_G(r)`

Definition and examples of relations

In the terminology of relational databases:

- A relation `r: A <=> B` is a table with 2 columns (`A` and `B`)
- Each row `(a: A, b: B)` means that the value `a` is related to `b`

Mathematically speaking: a relation `r: A <=> B` is a subset $r \subset A \times B$

- We write `(a, b) in r` to mean $a \times b \in r$ where $a \in A$ and $b \in B$

Relations can be many-to-many while functions `A => B` are many-to-one
A function `f: A => B` can be also viewed as a relation `rel(f): A <=> B`

- Two values `a: A`, `b: B` are in `rel(f)` if `b == f(a)`
- `rel(identity: A => A)` defines an **identity** relation `id: A <=> A`

Example of a relation that can be many-to-many:
Given two functions `f: A => C`, `g: B => C`, define a "pullback" relation `pullback(f, g): A <=> B` as: `(a: A, b: B) in r` means `f(a) == g(b)`

- The pullback relation is *not* equivalent to a function `A => B` or `B => A`

Proof of relational parametricity. I. Relation combinators

Relation combinators:

- For any relation `r: A <=> B`, the **inverse** relation is `inv(r): B <=> A`
  - The inverse operation is its own inverse: `inv(inv(r)) == r`
- For any relations `r: A <=> B` and `s: A <=> B`, get the union (`r or s`) and the intersection (`r and s`):

  `(a, b) in (r and s)` means `(a, b) in r` and `(a, b) in s`
  `(a, b) in (r or s)` means `(a, b) in r` or `(a, b) in s`
- For any relations `r: A <=> B` and `s: B <=> C` define the **composition** (`r compose s`) as a relation `u: A <=> C` by `(a: A, c: C) in u` when there exists `b: B` such that `(a, b) in r` and `(b, c) in s`
  - Composition corresponds to "join" in relational databases
  - Directionality law: `inv(r compose s) == inv(s) compose inv(r)`
  - Associativity and identity laws with respect to `id: A <=> A`
  - Preserves composition of functions: for `f: A => B` and `g: B => C`, `rel(f andThen g) == rel(f) compose rel(g)`
- The "pullback relation" can be expressed through composition:

  `pullback(f, g) == rel(f) compose inv(rel(g))`

Pullback relation expressed through composition of relations

For any `f: A => C`, `g: B => C`, `a: A`, `b: B`, to prove:

- `(a, b) in pullback(f, g)` is equivalent to:
  `(a, b) in rel(f) compose inv(rel(g))`

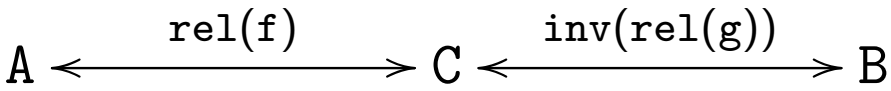

- The first condition is equivalent to: `f(a) == g(b)`
- The second condition is equivalent to: there exists `c: C` such that:
  `(a, c) in rel(f)` and `(c, b) in inv(rel(g))`
- This is equivalent to: `z` is such that `c == f(a)` and `c == g(b)`
- This is equivalent to the first condition

Proof of relational parametricity. II. Definition of `rmap`

For a type constructor `F` and `r: A <=> B`, need `rmap(r): F[A] <=> F[B]`

Define `rmap` for `F[A]` by induction over the *type expression* of `F[A]`

There are seven possibilities (assuming that the code is fully parametric):

1. `F[A] = Unit` or another fixed type (say, `T`) not related to `A`
2. The identity functor: `F[A] = A`
3. Product type: `F[A] = (G[A], H[A])`
4. Co-product type: `F[A] = Either[G[A], H[A]]`
5. Function type: `F[A] = G[A] => H[A]`
6. Recursive type: `F[A] = G[A, F[A]]`
7. Universally quantified term: `F[A] = [Z] => G[A, Z]`

Define `rmap` similarly to how a functor's `fmap` is defined in these cases

- The inductive assumption is that liftings to `G` and `H` are already defined
- For `G[A, Z]`, need to use two liftings (`rmap_G1` and `rmap_G2`)
- Liftings with respect to different type parameters will commute!
- For `F[A] = G[H[A]]` we expect `rmap_F(r) == rmap_G(rmap_H(r))`

Some diagrams for clarification

The commutativity theorem for relational liftings: For any type constructor `G[A, X]` and any two relations `r: A <=> B` and `s: X <=> Y`:

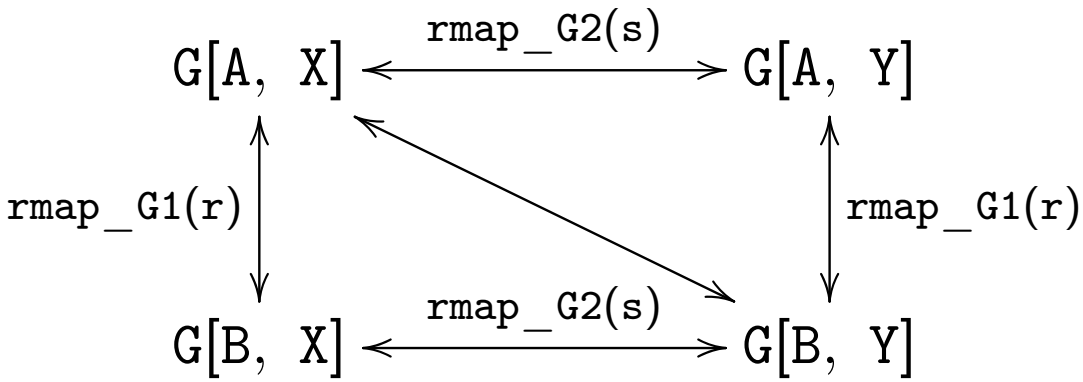

Relational lifting for a composition of type constructors, `F[A] = G[H[A]]`:

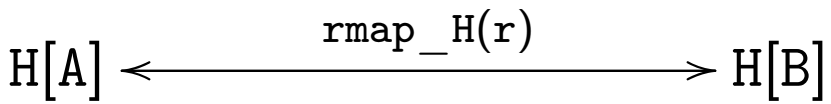

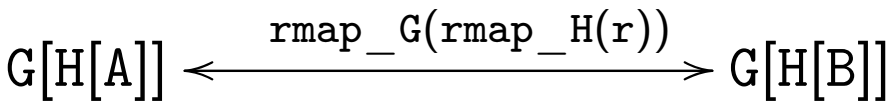

Proof of relational parametricity. II. Definition of `rmap`

Need to define `rmap(r): F[A] <=> F[B]` in the 7 cases:

1. `F[A] = T` (a fixed type): define `rmap(r) = id:  T <=> T`
2. The identity functor, `F[A] = A`: define `rmap(r) = r:  A <=> B`
3. When `F[A] = (G[A], H[A])`: define `((g1,h1), (g2,h2)) in rmap(r)` to mean `(g1, g2) in rmap_G(r)` and `(h1, h2) in rmap_H(r)`
4. When `F[A] = Either[G[A], H[A]]`: either `(Left(g1), Left(g2)) in rmap(r)` when `(g1, g2) in rmap_G(r)` or `(Right(h1), Right(h2)) in rmap(r)` when `(h1, h2) in rmap_H(r)`
5. When `F[A] = G[A] => H[A]`: define `(f1, f2) in rmap(r)` to mean `(f1(g1), f2(g2)) in rmap_H(r)` for any `g1: G[A]` and `g2: G[B]` such that `(g1, g2) in rmap_G(r)`
6. When `F[A] = G[A, F[A]]`: define `rmap(r) = rmap_G1(r) compose rmap_G2(rmap(r))` – the second `rmap(r)` is a recursive call
7. When `F[A] = [Z] => G[A, Z]`: define `(f1, f2) in rmap(r)` to mean: for any types `Z1` and `Z2`, and for any relation `s: Z1 <=> Z2`, we require `(f1[A][Z1], f2[B][Z2]) in (rmap_G1(r) compose rmap_G2(s))`

Proof of relational parametricity. III. Examples of using `rmap`

Use `rmap` to lift a relation `r` to a type constructor

Two main examples of relations generated by functions:

`rel(f)` and `pullback(f, g)`

Three main examples of type constructors (`F[A]`, `G[A]`, `H[A]`):

- If `F[A]` is covariant then:
  `rmap(rel(f)) == rel(fmap(f))`
  `rmap(pullback(f, g)) == pullback(fmap(f), fmap(g))`
- If `G[A] = A => A` then `(fa, fb) in rmap(rel(f))` means:
  when `(a, b) in rel(f)` then `(fa(a), fb(b)) in rel(f)`
  or: `f(fa(a)) == fb(f(a))` or: `fa andThen f == f andThen fb`
  This relation has the form of a pullback
- If `H[A] = (A => A) => A` then `(fa, fb) in rmap_H(rel(f))` means:
  when `(p, q) in rmap_G(rel(f))` then `(fa(p), fb(q)) in rel(f)`
  equivalently: if `p andThen f == f andThen q` then `f(fa(p))==fb(q)`
  This is *not* a pullback relation: cannot express `p` through `q`

It is hard to use relations that do not have the form of a pullback

Proof of relational parametricity. IV. Formulation

Instead of proving relational properties for `t[A]: P[A] => Q[A]`, use the function type and the quantified type constructions and get:

- Any fully parametric `t[A]: P[A]` satisfies for any `r: A <=> B` the relation `(t[A], t[B]) in rmap_P(r)`
- Any fully parametric `t: P[]` satisfies `(t, t) in rmap_P(id)`

It is more convenient to prove a parametricity theorem with a free variable:

- Any fully parametric expression `t[A](z): P[A]` with `z: Q[A]` satisfies, for any relation `r: A <=> B` and for any `z1: Q[A]`, `z2: Q[B]`, the law: if `(z1, z2) in rmap_Q(r)` then `(t[A](z1), t[B](z2)) in rmap_P(r)`

This applies to expressions containing one free variable (`z`)

- Any number of free variables can be grouped into a tuple

Proof of relational parametricity. V. Outline

The theorem says that `t[A](z)` satisfies its relational parametricity law
Proof goes by induction on the structure of the code of `t[A](z)`
At the top level, `t[A](z)` must have one of the 9 code constructions
Each construction decomposes the code of `t[A](z)` into sub-expressions
The inductive assumption is that the theorem holds for all sub-expressions (including the bound variable `z`)

Proof of relational parametricity. VI. Examples

We will show how to prove the first 4 constructions

Constant type: If `t[A](z) = c` where `c` is a fixed value of a fixed type `C`:

- We have `rmap_P(r) == id` while `(c, c) in id` holds

Use argument: If `t[A](z) = z` where `z` is a value of type `Q[A]`:

- If `(z1, z2) in rmap_Q(r)` then `(t(z1), t(z2)) in rmap_Q(r)`

Create function: If `t(z) = h => s(z, h)` where `h: H[A]` and `s(z, h): S[A]`:

- If `(z1, z2) in rmap_Q(r)` and `(h1, h2) in rmap_H(r)` then `(s(z1, h1), s(z2, h2)) in rmap_S(r)`

Use function: If `t(z) = g(z)(h(z))` where `g(z): H[A] => P[A]` and `h(z): H[A]` are sub-expressions:

- If `(z1, z2) in rmap_Q(r)` then inductive assumption says: `(h(z1), h(z2)) in rmap_H(r)`
- If `(h1, h2) in rmap_H(r)` then inductive assumption says: `(g(h1), g(h2)) in rmap_P(r)`

Summary

- Relational parametricity is a powerful technique
- It has been generalized to many different settings
  - Gradual typing, higher-kinded types, dependent types, etc.
- Relational parametricity has a steep learning curve
  - Cannot directly write code that manipulates relations
  - All calculations need to be done symbolically or with proof assistants
- The result may be a relation that is difficult to interpret as code
- A couple of results in FP do require the relational naturality law
- More details in the free book — https://github.com/winitzki/sofp